

# COMP SCI 1400

## AI Technologies

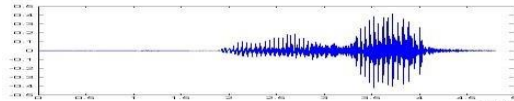
Dr. Kamal Mammadov

# Outline

- Classification
  - Support Vector Machine
- Machine Learning
  - ✓ supervised learning
  - ✓ unsupervised learning
  - semi-supervised learning
  - reinforcement learning

# Machine Learning --- Finding Functions

- Speech Recognition

$$f(\text{  }) = \text{"How are you"}$$

- Image Recognition

$$f(\text{  }) = \text{"Cat"}$$

- Playing Go

$$f(\text{  }) = \text{"5-5"} \text{ (next move)}$$


- Dialogue System

$$f(\text{ "How are you?" } \text{ (what the user said) }) = \text{ "I am fine." } \text{ (system response)}$$

# Supervised Learning

## Unsupervised Learning

# Supervised Learning



$x_1$ :



$y_1$ : "Cat"

$x_2$ :



$y_2$ : "Cat"

$x_3$ :



$y_3$ : "Dog"

$x_4$ :



$y_4$ : "Dog"

Labelled Data



# Support Vector Machine

- Support Vector Machine --- 1D



The **red dots** represent mice are **not obese**...


- Support Vector Machine --- 1D



...and the **green dots** represent mice are **obese**.




- Support Vector Machine --- 1D




Based on these observations, we can pick a threshold...

- Support Vector Machine --- 1D




...and when we get a new observation that  
has less mass than the threshold...

- Support Vector Machine --- 1D




And when we get a new observation with more mass than the threshold...

- Support Vector Machine --- 1D




However, what if get a new observation here?

- Support Vector Machine --- 1D



Because this observation has more mass than the threshold, we classify it as **obese**.


- Support Vector Machine --- 1D




But that doesn't make sense, because it is much closer to the observations that are *not obese*.

- Support Vector Machine --- 1D

Find a new threshold




- Support Vector Machine --- 1D






- Support Vector Machine --- 1D




- Support Vector Machine --- 1D




So it makes sense to classify this new observation as *not obese*.

- Support Vector Machine --- 1D



- Support Vector Machine --- 1D




When the threshold is halfway between the two observations, the **margin** is as large as it can be.


We maximize the margin to both classes to find the best classifier

Support vector determines the classifier


- Support Vector Machine --- 1D




- Support Vector Machine --- 1D



- Support Vector Machine --- 1D




- Support Vector Machine --- 1D




In this case, the **Maximum Margin Classifier** would be super close to the *obese* observations...



- Support Vector Machine --- 1D




- Support Vector Machine --- 1D



...we would classify it as **not obese**, even though most of the **not obese** observations are much further away than the **obese** observations.

- Support Vector Machine --- 1D



So **Maximal Margin Classifiers** are super sensitive to outliers in the training data and that makes them pretty lame.

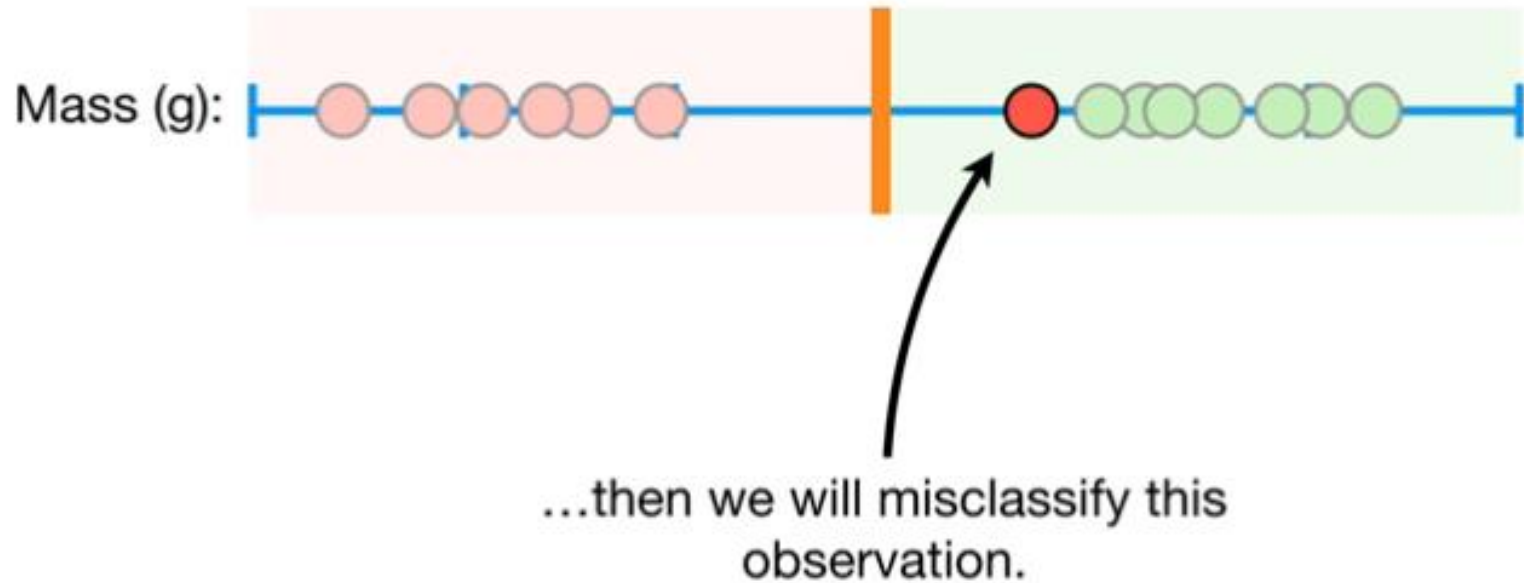
- Support Vector Machine --- 1D

## Solution




To make a threshold that is not so sensitive to outliers we must **allow misclassifications**.


- Support Vector Machine --- 1D



- Support Vector Machine --- 1D




- Support Vector Machine --- 1D



Choosing a threshold that allows misclassifications is an example of the **Bias/Variance Tradeoff** that plagues all of machine learning.


- Support Vector Machine --- 1D



In other words, before we allowed misclassifications, we picked a threshold that was very sensitive to the training data (low bias)...



- Support Vector Machine --- 1D




Overfitting

...and it performed poorly when  
we got new data (high variance).


- Support Vector Machine --- 1D

## Solution



When we allow misclassifications, the distance between the observations and the threshold is called a **Soft Margin**.


- Support Vector Machine --- 1D



When we use a **Soft Margin** to determine the location of a threshold


then we are using a **Soft Margin Classifier** aka  
a **Support Vector Classifier** to classify  
observations.

- Support Vector Machine --- 1D



How do we know which soft margin is better?


- Support Vector Machine --- 1D



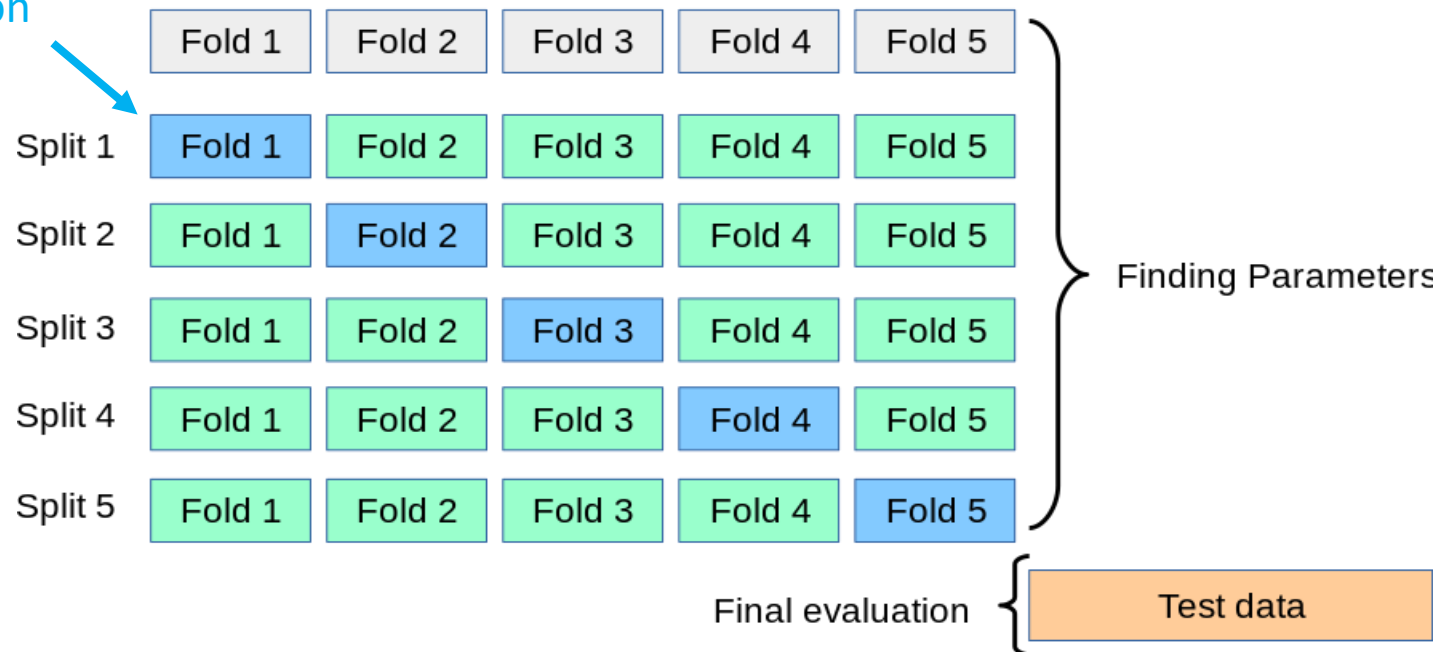
The answer is simple: We use **Cross Validation** to determine how many misclassifications and observations to allow inside of the **Soft Margin** to get the best classification.

# Support Vector Machine


- K-fold cross validation
- $k=5$




Validation



- Support Vector Machine --- 2D




- Support Vector Machine --- 2D







- Support Vector Machine --- 2D




- Support Vector Machine --- 2D



- Support Vector Machine --- 3D



- Support Vector Machine --- 3D




- Support Vector Machine --- >3D

**NOTE:** If we measured *mass*,  
*height*, *age* and *blood pressure*,  
then the data would be in **4**  
**Dimensions...**

The Support Vector Classifier is a **hyperplane**

# How to find the optimal classifier?



<https://www.pycodemates.com/2022/09/primal-formulation-of-svm-simplified.html>

[https://en.wikipedia.org/wiki/Support\\_vector\\_machine](https://en.wikipedia.org/wiki/Support_vector_machine)

# Exercise/Homework: Optimal Maximal Margin Classifier

Labelled dataset  $(x_1, y_1), (x_2, y_2), \dots, (x_d, y_d)$ , where  $x_i \in \mathbb{R}^n$ ,  $y_i \in \{-1, 1\}$  for all  $i$ .

We need to find parameters  $w \in \mathbb{R}^n$ ,  $b \in \mathbb{R}$ , satisfying for all  $i$ :


$$y_i = \begin{cases} 1 & w^T x_i + b \geq 1 \\ -1 & w^T x_i + b \leq -1 \end{cases}$$

That maximizes the distance between the two hyperplanes.

Question: Given any  $w$  and  $b$ , how do we compute the distance between the hyperplanes?

Hint: First prove that vector  $w$  is orthogonal to any vector in the hyperplane.

# Optimal Soft Margin Classifier




<https://www.pycodemates.com/2022/09/primal-formulation-of-svm-simplified.html>

[https://en.wikipedia.org/wiki/Support\\_vector\\_machine](https://en.wikipedia.org/wiki/Support_vector_machine)




- Support Vector Machine



How can we find a classifier that can tell whether a new dosage can cure or not?

- Support Vector Machine

Can we do better than **Maximal Margin Classifiers** and **Support Vector Classifiers**?




- Support Vector Machine


# Support Vector Machines!!!




- Support Vector Machine




- Support Vector Machine




- Support Vector Machine



- Support Vector Machine




- Support Vector Machine







- Support Vector Machine




- Support Vector Machine




- Support Vector Machine




- Support Vector Machine



- Support Vector Machine



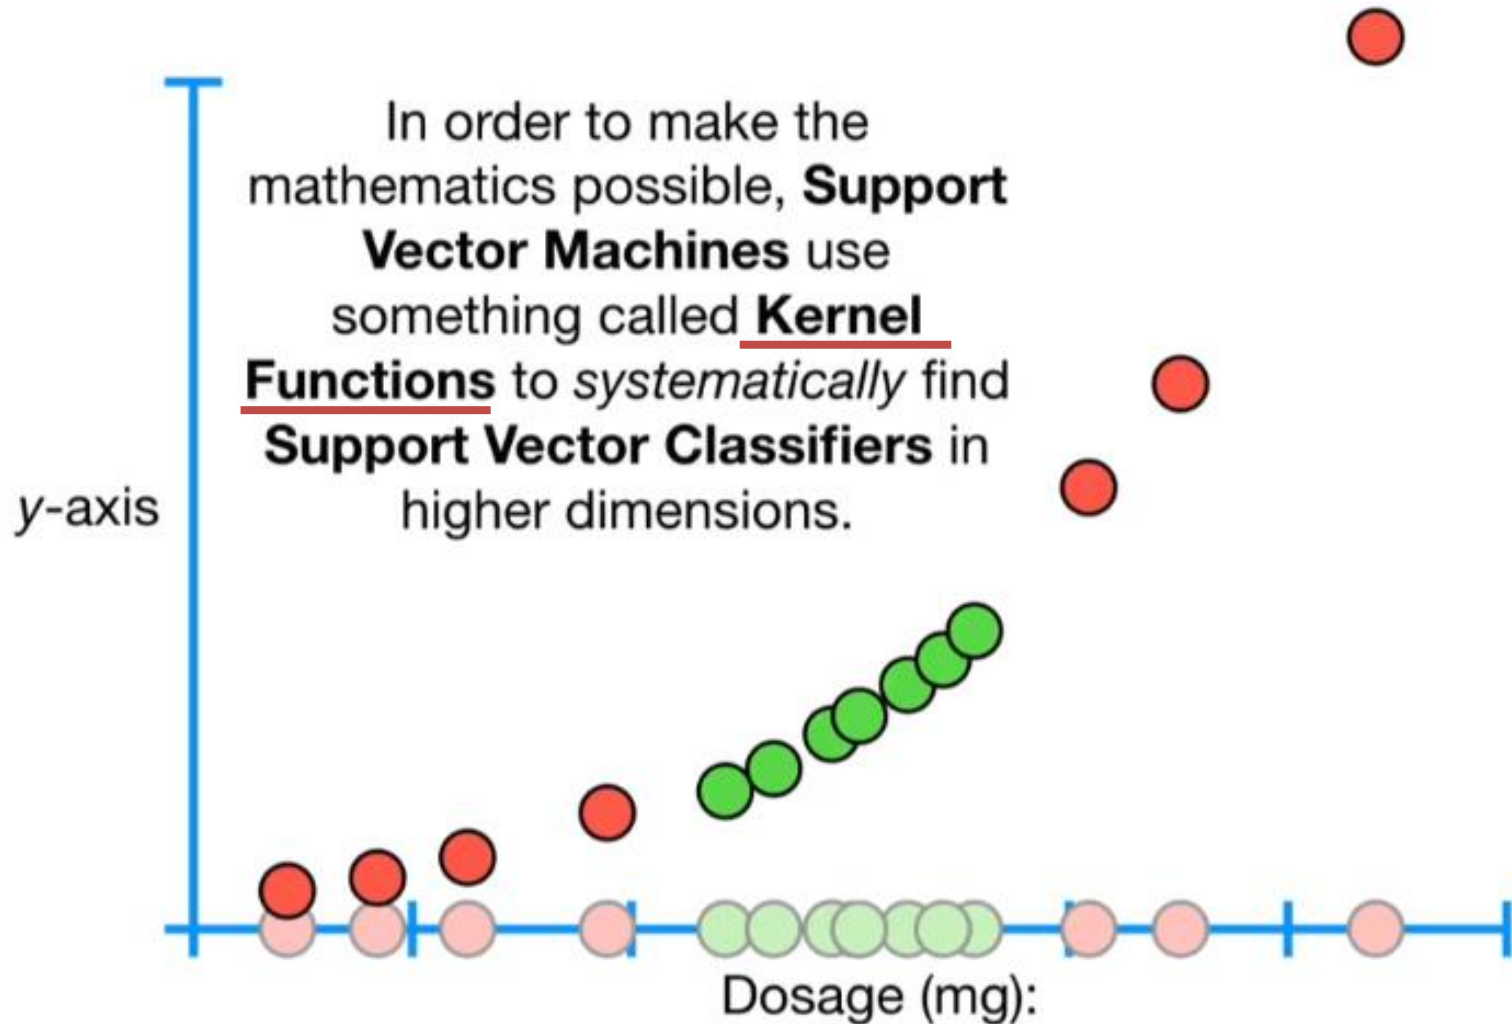
- Support Vector Machine



- Support Vector Machine

How do we know what kind of higher dimensions  
should be used?

- Support Vector Machine





- Support Vector Machine

Polynomial kernel


$$k(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i \cdot \mathbf{x}_j)^d$$

Gaussian kernel


$$k(\mathbf{x}_i, \mathbf{x}_j) = \exp\left(-\gamma\|\mathbf{x}_i - \mathbf{x}_j\|^2\right)$$

Using these kernel techniques, we can get classifiers that work in higher dimensions

- Support Vector Machine



- Support Vector Machine



- Support Vector Machine

How do we find the best hyper-parameters?

- Support Vector Machine

1. We train models with different hyper-parameters
2. Compare their performance

# Model Evaluation

$$\min_{\mathbf{w}, \mathbf{b}, \xi} \frac{1}{2} \|\mathbf{w}\|^2 + C \sum_{i=1}^n \xi_i$$

Example: SVC with Polynomial kernel

$$k(\mathbf{x}_i, \mathbf{x}_j) = (\mathbf{x}_i \cdot \mathbf{x}_j)^d$$

Setting A,  $d=2$ ,  $C=1$

$$\hat{y} = f_A(x),$$

Setting B,  $d=2$ ,  $C=10$

$$\hat{y} = f_B(x),$$

Say we have 145 mice for testing, we get two groups of predictions

Model A:  $f_A(x_1), f_A(x_2), \dots, f_A(x_{145})$

Model B:  $f_B(x_1), f_B(x_2), \dots, f_B(x_{145})$

- Evaluation

Model A:  $f_A(x_1), f_A(x_2), \dots, f_A(x_{145})$

		Ground truth	
		Cured	Not Cured
Prediction	Cured	True Positive	False Positive
	Not Cured	False Negative	True Negative

$$\begin{aligned}\text{Accuracy} &= \frac{\text{Correct Prediction}}{\text{Total Prediction}} \\ &= \frac{\text{True Positive} + \text{True Negative}}{\text{True Positive} + \text{False Positive} + \text{True Negative} + \text{False Negative}}\end{aligned}$$

- Evaluation

Results of model A, under split1 setting

		Ground truth	
		Cured	Not Cured
Prediction	Cured	True Positive 34	False Positive 23
	Not Cured	False Negative 12	True Negative 76

$$\begin{aligned}\text{Accuracy} &= \frac{\text{True Positive} + \text{True Negative}}{\text{True Positive} + \text{False Positive} + \text{True Negative} + \text{False Negative}} \\ &= \frac{34+76}{34+23+76+12}\end{aligned}$$



- Evaluation

How do we know the performance of a model is good or bad?

		Ground truth	
		Cured	Not Cured
Prediction	Cured	True Positive 34	False Positive 23
	Not Cured	False Negative 12	True Negative 76

Sensitivity or  
True Positive Rate or  
Recall

$$\begin{aligned} &= \frac{\text{True Positive}}{\text{Actual Positive}} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Negative}} \\ &= \frac{34}{34+12} \end{aligned}$$

- Evaluation

How do we know the performance of a model is good or bad?

		Ground truth	
		Cured	Not Cured
Prediction	Cured	True Positive 34	False Positive 23
	Not Cured	False Negative 12	True Negative 76

$$\begin{aligned}\text{Precision} &= \frac{\text{True Positive}}{\text{Predicted Positive}} = \frac{\text{True Positive}}{\text{True Positive} + \text{False Positive}} \\ &= \frac{34}{34+23}\end{aligned}$$

- Evaluation


How do we know the performance of a model is good or bad?

		Ground truth	
		Cured	Not Cured
Prediction	Cured	True Positive 34	False Positive 23
	Not Cured	False Negative 12	True Negative 76

$$\begin{aligned}\text{False Positive Rate} &= \frac{\text{False Positive}}{\text{Actual Negative}} = \frac{\text{False Positive}}{\text{False Positive} + \text{True Negative}} \\ &= \frac{23}{23+76}\end{aligned}$$


- # Evaluation

Receiver Operating Characteristic curve (ROC curve)



Closer to left upper corner, better performance

# ROC curve



# Theorem (Random classifier)

A random classifier lies anywhere on the line  $TPR = FPR$ .


Proof: Let  $n$  be the number of participants/test subjects that our classifier is tasked with predicting whether they are cured. Let  $p$  denote the probability that a person is cured.

Given a random classifier, that randomly predicts a person is cured with probability  $q$  and predicts a person is not cured with probability  $1 - q$ .

Prove that  $E[TPR] = E[FPR]$  holds.

Hint: First find the values for  $E[TP]$ ,  $E[FP]$ ,  $E[FN]$ ,  $E[TN]$ ; then use the formulas  $FPR = \frac{FP}{FP+TN}$ ,  $TPR = \frac{TP}{TP+FN}$ .

# SVM Example



# SVM Example

conda install python=3.10

conda install matplotlib

conda install scikit-learn

```
import matplotlib.pyplot as plt
```

```
from sklearn import svm
```

```
from sklearn.datasets import make_blobs
```

```
from sklearn.inspection import DecisionBoundaryDisplay
```

```
# we create 40 separable points
```


```
X, y = make_blobs(n_samples=40, centers=2, random_state=6)
```

```
plt.scatter(X[:, 0], X[:, 1], c=y, s=30, cmap=plt.cm.Paired)
```

```
# fit the model, don't regularize for illustration purposes
```

```
clf = svm.SVC(kernel="linear", C=1000)
```

```
clf.fit(X, y)
```







# Example

```
# plot the decision function
ax = plt.gca()
DecisionBoundaryDisplay.from_estimator(
    clf,
    X,
    plot_method="contour",
    colors="k",
    levels=[-1, 0, 1],
    alpha=0.5,
    linestyles=["--", "-", "--"],
    ax=ax,
)
# plot support vectors
ax.scatter(
    clf.support_vectors_[:, 0],
    clf.support_vectors_[:, 1],
    s=100,
    linewidth=1,
    facecolors="none",
    edgecolors="k",
)
plt.show()
```

Can SVM do multiclass classification?

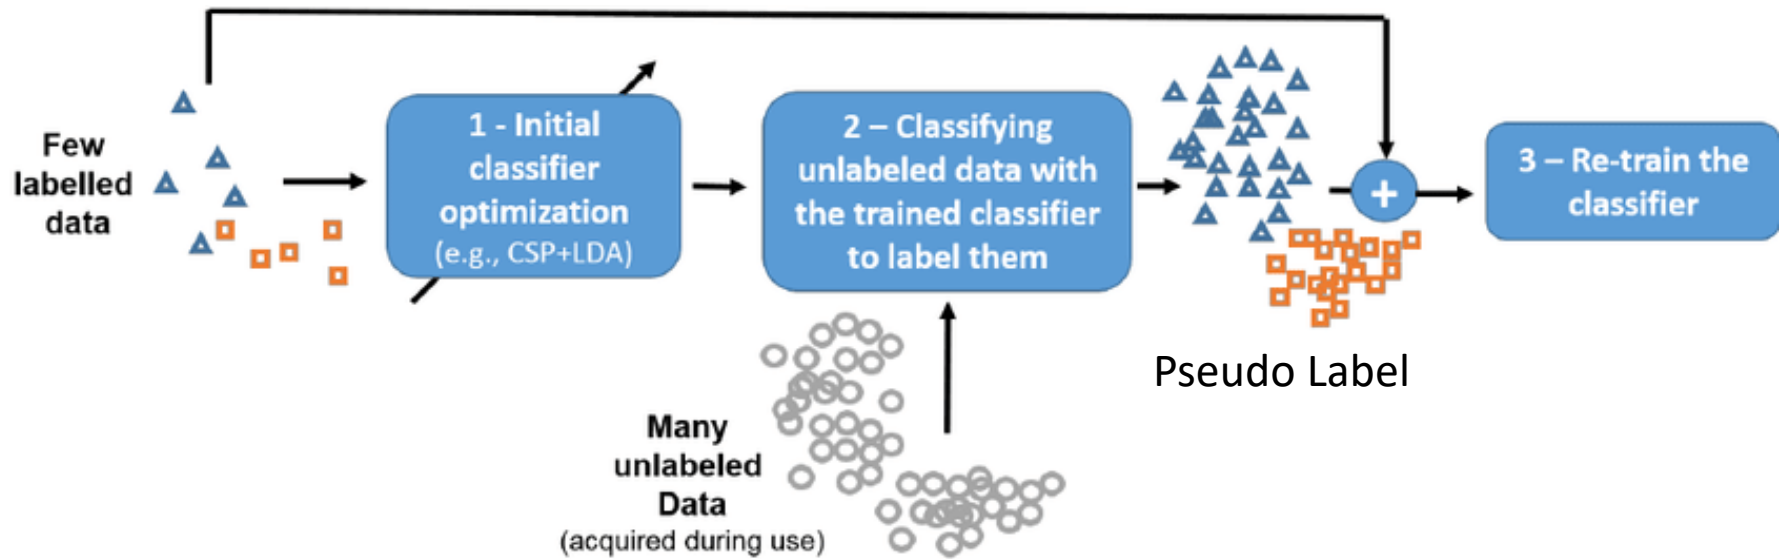


Can SVM do multiclass classification?



# Semi-supervised Learning

# Semi-supervised Learning



- Acknowledgement

Part of the material are from Josh Starmer

<https://www.youtube.com/watch?v=efR1C6CvhmE>